

Cross compilation with Clang and LLVM tools

Peter Smith Linaro TCWG



Linaro
connect

Bangkok 2019

Contents

- Refresher on cross compilation.
- Clang/LLVM toolchains.
- Using Clang as a cross compiler.
- Additional cross compilation options.
- Using CMake.
- Differences between GCC and Clang.
- Assembling a Clang toolchain today.

Definitions

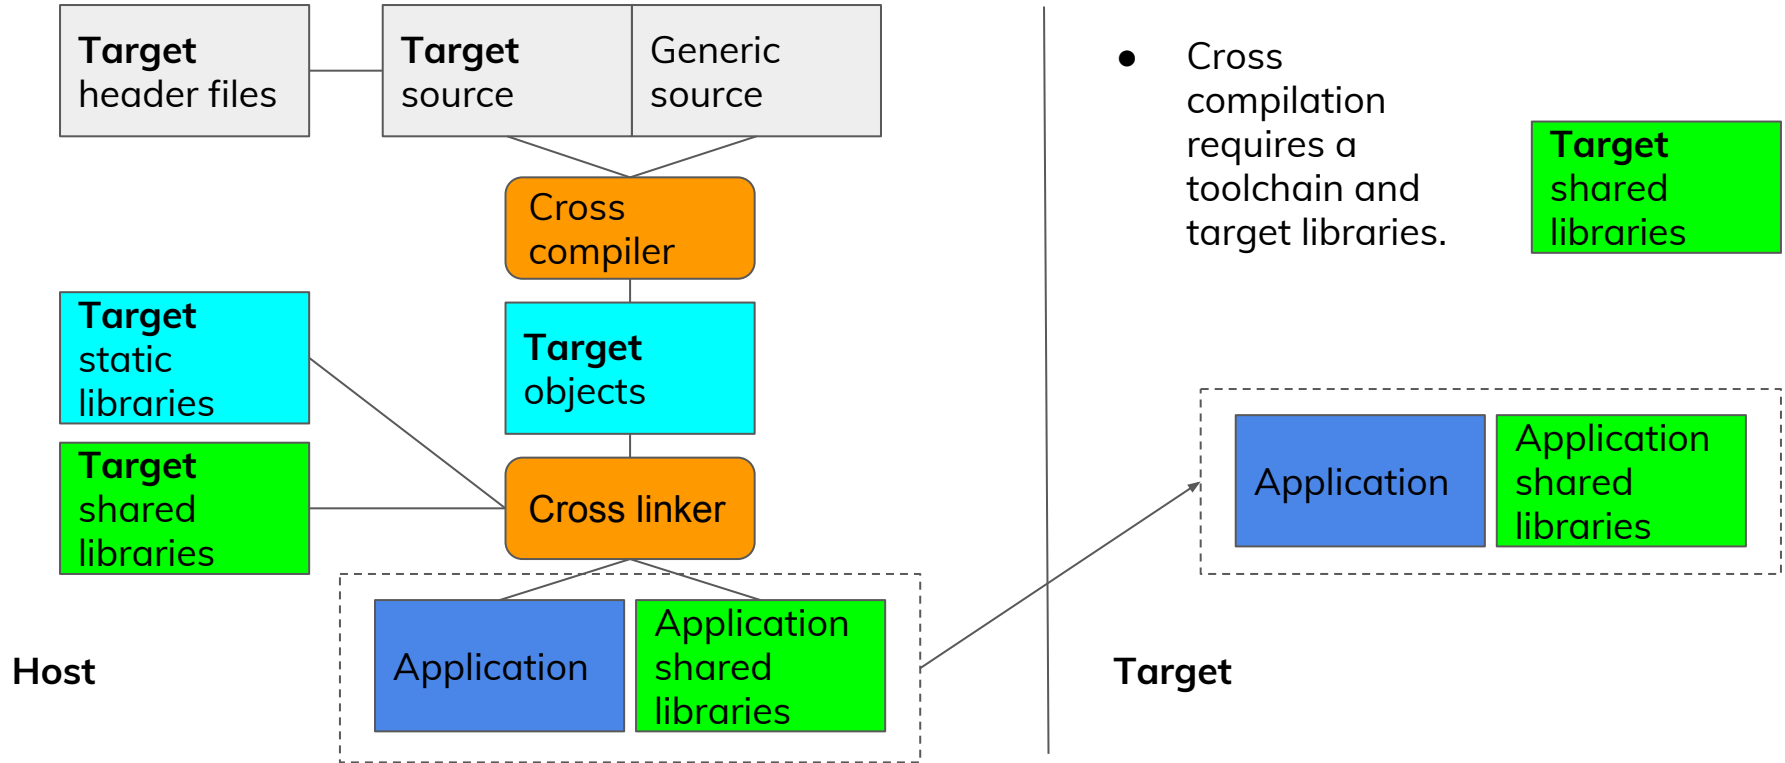
- **Host**
 - The system that we run the development tool on.
- **Target**
 - The system that we run the generated program on.
- **Native** compilation
 - Host is the same as Target.
- **Cross** compilation
 - Host is different to Target.

Motivation for cross compilation

- Productivity boost when host is faster than target.
- Target can't run a C/C++ compiler.
- Need to build program for many architectures.
- Bootstrapping a compiler on a new architecture.



A cross compiled application



Clang/LLVM Toolchains

- What is provided in an installation?
- What is missing?

Cross compilation and the LLVM toolchain

- Clang and other LLVM tools can work with multiple targets from the same host binary.
- Clang and LLD drivers can emulate the drivers of other toolchains.
- Controlled by the target triple.
- LLVM project does not have implementations of all the parts of toolchain.
- LLVM project includes some but not all of the library dependencies.

Toolchain components

Component	LLVM	GNU
C/C++ Compiler	clang	gcc
Assembler	clang integrated assembler	as
Linker	ld.lld	ld.bfd, ld.gold
Runtime	compiler-rt	libgcc
Unwinder	libunwind	libgcc_s
C++ library	libc++abi, libc++	libsupc++, libstdc++
Utils such as archiver	llvm-ar, llvm-objdump etc.	ar, objdump etc.
C library		glibc, newlib

Toolchain components used by Clang

- Defaults chosen at build time, usually favour GNU libraries for Linux targets, otherwise LLVM.
- Compiler runtime library
 - `--rtlib=compiler-rt`, `--rtlib=libgcc`.
 - `compiler-rt` needed for sanitizers but these are separate from builtins provided by `libgcc`.
- C++ library
 - `--stdlib=libc++`, `--stdlib=libstdc++`.
 - No run-time option to choose C++ ABI library, determined at C++ library build time.
- Linker
 - `-fuse-ld=lld`, `-fuse-ld=bfd`, `-fuse-ld=gold`.
 - Driver calls `ld.lld`, `ld.bfd`, `ld.gold` respectively.
- C-library choice can affect target triple
 - For example `arm-linux-gnueabi`, `arm-linux-musleabi`.

Using Clang as a Cross Compiler

- Deconstructing the Target Triple.
- Using a GCC toolchain to provide missing components.
- The Clang driver.
- How a Clang installation is laid out.

Target Triple

- General format of **<Arch><Sub-arch>-<Vendor>-<OS>-<Environment>**
 - **Arch** is the architecture that you want to compile code for
 - Examples include arm, aarch64, x86_64, mips.
 - **Sub-arch** is a refinement specific to an architecture
 - Examples include armv7a armv7m.
 - **Vendor** captures differences between toolchain vendors
 - Examples include Apple, PC, IBM.
 - **OS** is the target operating system
 - Examples include Darwin, Linux, OpenBSD, none.
 - **Environment** includes the ABI and object file format
 - Examples include android, elf, gnu, gnueabihf.
- Missing parts replaced with “unknown”.

Arm Target Triple

- **<Arch>**
 - arm or thumb. The -marm and -mthumb take precedence.
- **<Sub-arch>**
 - Accepts several forms, v7a, v7-a. -march, -mcpu take precedence.
- **<OS>**
 - linux, android, none (for bare metal)
- **<Environment>**
 - gnueabi, gnueabihf, eabi, eabihf, musleabi, musleabihf
 - hf is for hard-float, -mfloat-abi takes precedence.

Arm Target Triple examples


- **thumbv8-m.mainline-none-eabi**
 - bare metal soft-floating point Armv8-m.mainline.
 - For M-class, thumb is always selected even if arm is in the triple.
- **armv7a-linux-gnueabi**
 - Linux, hard-floating point Armv7-a)
- **arm-linux-gnueabi -march=armv7-a -mfloat-abi=softfp**
 - Linux, soft-floating calling convention, Armv7-a

Clang driver and toolchains

- Driver mode
 - gcc, g++, cpp (preprocessor), cl (MSVC).
 - Set with option or inferred from filename clang, clang++, clang-cl.
- Target triple used to instantiate a ToolChain derived class
 - arm-linux-gnueabihf instantiates the Linux ToolChain (Linux.cpp, Gnu.cpp).
 - arm-none-eabi instantiates the bare metal ToolChain (BareMetal.cpp).
- Toolchain class has knowledge of how to emulate the native toolchain
 - Include file locations.
 - Library locations.
 - Constructing linker and non integrated assembler options.
 - Includes cross compilation emulation.
- Not all functionality is, or could realistically be, documented.

Clang Driver

```
clang --target=aarch64-linux-gnu func1.s hello.c -o hello
```



Architecture: aarch64
Sub-architecture: not applicable
Vendor unknown
OS: linux
Environment: GNU

```
clang-7.0 -cc1 -triple  
aarch64-linux-gnu  
-target-cpu=generic  
-target-feature +neon  
-target-abi aapcs  
-Isystem /path/to/includes  
...
```

```
clang-7.0 -cc1as -triple  
aarch64-linux-gnu  
-target-cpu=generic  
-target-feature +neon  
...
```

```
ld.lld -m aarch64linux  
-dynamiclinker  
/lib/ld-linux-aarch64.so  
-L /path/to/system/libraries  
-lc  
...
```


Anatomy of a Clang installation

- clang+llvm-7.0.1-x86_64.linux-gnu-ubuntu-16.04
 - bin (host user and developer tools such as clang, lld, llvm-mc)
 - include
 - c++ (libc++)
 - clang, clang-c, llvm, llvm-c (for compiling against llvm as a library)
 - lib
 - libc++, libc++abi, libgomp, (for user programs, **not namespaced by target**)
 - libclang, libllvm, (for linking against llvm as a library)
 - clang
- resource dir → ● 7.0.0
 - include
 - stdint.h, arm_aacle.h, arm_neon.h and other compiler specific headers
 - sanitizers, xray subdirs
 - lib
 - linux (for OS=linux, can also be baremetal)
- compiler-rt → ● libclang_rt.builtins-x86_64.a (and other compiler-rt libs)
 - share
 - asan_blacklist.txt (and other sanitizer blacklists)
- share scripts, man pages etc.

Using a Clang installation for cross compiling

- All the host tools will be able to cross compile for Arm
 - Assuming toolchain builder didn't exclude it from build.
- The libc++, libc++abi and all the LLVM/Clang libraries will be for the host.
 - Cannot use libc++ with the default library paths.
 - Linux defaults to libstdc++
- The compiler-rt library will also be for the host
 - It has the target in the name so it is possible to install Arm versions
 - Linux defaults to libgcc
- The remainder of the libraries will need to be found outside the installation
 - Linux distributions multiarch.
 - A separate GCC installation.

Finding the location of a GCC cross compiler

- A bunch of heuristics guided by two relevant command-line options
 - `--gcc-toolchain`
 - `--sysroot`
- Clang looks for `lib/gcc/<gcc-triple>/major.minor.patch` and `lib/gcc-cross/<gcc-triple>/major.minor.patch`
 - Prefixed first with `--gcc-toolchain`, then `--sysroot`.
 - The highest GCC version found is chosen.
 - GCC toolchain components assumed to be relative to this location.
- Includes and libraries searched for in `sysroot/usr`
- The Arm and Linaro toolchains need both `--gcc-toolchain` and `--sysroot`
 - `--gcc-toolchain=/path/to/install-dir`
 - `--sysroot=/path/to/install-dir/<gcc-triple>/libc`

Example multiarch on Ubuntu 16.04

```
$ clang hello.c --target=arm-linux-gnueabihf -o hello
```

```
$ qemu-arm -L /usr/arm-linux-gnueabihf hello
```

Hello World

```
$ clang hello.c -v --target=arm-linux-gnueabihf -o hello
```

```
Found candidate GCC installation: /usr/lib/gcc-cross/arm-linux-gnueabihf/5.4.0
```

```
Selected GCC installation: /usr/lib/gcc-cross/arm-linux-gnueabihf/5.4.0
```

```
ignoring nonexistent directory "/include"
```

```
#include "... " search starts here:
```

```
#include <...> search starts here:
```

```
/usr/local/include
```

```
/path/to/clang/lib/clang/9.0.0/include
```

```
/usr/include/arm-linux-gnueabihf
```

```
/usr/include
```

```
End of search list.
```

```
...
```

Potential for host
pollution

Example aarch64-linux-gnu toolchain

```
$ clang hello.c --target=aarch64-linux-gnu -o hello
--gcc-toolchain=/work/gcc7a64 --sysroot=/work/gcc7a64/aarch64-linux-gnu/libc
$ qemu-aarch64 -L /work/gcc7a64/aarch64-linux-gnu/libc hello
```

Hello World

```
clang hello.c --target=aarch64-linux-gnu -o hello -v
--gcc-toolchain=/work/gcc7a64 --sysroot=/work/gcc7a64/aarch64-linux-gnu/libc
Found candidate GCC installation: /work/gcc7a64/lib/gcc/aarch64-linux-gnu/7.1.1
Selected GCC installation: /work/gcc7a64/lib/gcc/aarch64-linux-gnu/7.1.1
ignoring nonexistent directory "/work/gcc7a64/aarch64-linux-gnu/libc/usr/local/include"
ignoring nonexistent directory "/work/gcc7a64/aarch64-linux-gnu/libc/include"
#include "... " search starts here:
#include <...> search starts here:
    /path/to/clang/build/lib/clang/9.0.0/include
    /work/gcc7a64/aarch64-linux-gnu/libc/usr/include
End of search list.
```

Limitations of Clang's Driver

- No support for specs files
 - Clang has configuration files that can be used to partially emulate.
- No support for Linux or bare metal multilib
 - Android multilib support is available.
- Heuristics to find GCC installation are opaque and incomplete
 - Will work with well known distros.
- With multiarch there is a danger of host include pollution
 - `/usr/local/include` and `/usr/include` on constructed include path.
- Not easy to use libc++ when using `--gcc-toolchain`, `--sysroot`
 - These are in Clang's include and lib dir, not GCC's.
 - Need to either specify include and library dirs, or copy libc++ to GCC installation.

Additional cross compilation options

- Using configuration files.
- Using libc++ and compiler-rt.
- bare metal for embedded systems.


Using Clang configuration files

- Can be used to emulate some of the features of specs files
- `clang --config <config file>`
 - Contain command line options.
 - Can include other configuration files.
- Search path for config file can be set at build time
 - User dir: `-DCLANG_CONFIG_USER_DIR`
 - System dir: `-DCLANG_CONFIG_SYSTEM_DIR`
 - Directory where clang executable resides.
- Can select a config file automatically via renaming/symlinking clang
 - `armv7l-clang` is equivalent to `clang --config armv7l.cfg`
- No conditional features like in specs file
 - Would need a configuration file for each combination of choices.
 - For example `armv6-m-rdimon-nano-clang`
 - Could make individual config files that could be included.

Using libc++, compiler-rt

- So far we've been using a Clang installation with cross-tools but host LLVM libraries
 - Our libraries have come from a GCC cross toolchain.
 - Will need compiler-rt to use of sanitizers.
- Need to obtain target builds of LLVM libraries
 - Rebuild them from source.
 - Copy from a native Target Clang installation.
- Building from source is the only way to eliminate any dependencies on GNU libraries such as libgcc and libunwind.
- Using libc++ when cross compiling requires us to manually set the C++ include path to avoid conflicts with libstdc++.

Building Compiler-rt and Libc++ from source



1. Clang is needed for builtins and sanitizers, can skip if you already have a recent Clang.
2. The builtins are the equivalent of libgcc, a static library. These are a small part of compiler-rt that we need to build first.
3. The LLVM unwinder may use builtins from libunwind, needs C++ header from libc++
4. The libc++abi C++ runtime uses the unwinder from LLVM unwind
5. The libc++ C++ standard library uses libc++abi
6. The remainder of compiler-rt includes the sanitizers, which need libc++

A C++ example using libc++ and sanitizers

- Test is a slightly modified version of the ubsan example
 - Modified to throw an exception.
- We want to use as much of the LLVM libraries as possible
 - compiler-rt
 - libc++, libc++abi, libunwind

```
#include <iostream>
#include <string>

int func(void) {
    throw std::string("Hello World\n");
    return 0;
}

int main(int argc, char** argv) {
    try {
        func();
    } catch (std::string& str) {
        std::cout << str;
    }
    int k = 0x7fffffff;
    k += argc; //signed integer overflow
    return 0;
}
```

A C++ example using libc++ and sanitizers

```
$root=/path/to/clang/install_dir
$gcctoolchain=/path/to/gcc
$sysroot=${gcctoolchain}/aarch64-linux-gnu/libc
$ ${root}/bin/clang++ --target=aarch64-linux-gnu -fsanitize=undefined \
    --rtlib=compiler-rt --stdlib=libc++ \
    -nostdinc++ -I${root}/include/c++/v1 \
    -Wl,-L${root}/lib \
    --sysroot ${sysroot} \
    --gcc-toolchain=${gcctoolchain} \
    -rpath ${root}/lib \
    example.cpp -o example

$ qemu-aarch64 -L ${sysroot} example
Hello World
example.cpp:16:7: runtime error: signed integer overflow: 2147483647 + 1 cannot be
represented in type 'int'
```

Cross compiling for embedded systems

- Clang has a bare metal driver for Arm that is selected with `--target=arm-none-eabi`
- The functionality is somewhat bare
 - Setup include paths for libc++.
 - A kind of multiarch for the compiler-rt builtins.
 - Implicitly adding `-lc -lm`
 - Defaults to LLD, compiler-rt and the integrated assembler.
- Limitations
 - No support for finding an arm-none-eabi-gcc toolchain.
 - No support for specs files or multilib.
 - LLD support for linker scripts and embedded features good but not perfect.
- May need to use GNU supporting tools arm-none-eabi-gcc-objcopy
 - There are LLVM equivalents but they are not drop in replacements.
- A skeleton that you have to provide all the details yourself
 - Functional but not user-friendly.

Obtaining compiler-rt builtins

- Unless `--nostdlib` is used, Clang will add `-lclang_rt.builtins.<arch>.a`
- The only way to get armv6m, armv7m, armv8m versions is to build from source
- Guide available at <https://llvm.org/docs/HowToCrossCompileBuiltinsOnArm.html>
- Place built libraries into `lib/clang/7.0.0/lib/baremetal` directory

Using CMake

- Providing the target, sysroot and GCC toolchain
- Overriding the linker

Cross compilation with CMake

- Useful CMake options

- <https://cmake.org/cmake/help/v3.14/manual/cmake-toolchains.7.html>
- To skip link stage of trycompile stage
 - -DCMAKE_TRY_COMPILE_TARGET_TYPE=STATIC_LIBRARY
- Select Clang as compiler and assembler
 - -DCMAKE_C_COMPILER=clang
 - -DCMAKE_CXX_COMPILER=clang++
 - -DCMAKE_ASM_COMPILER=clang
- Set --sysroot with -DCMAKE_SYSROOT
- Set --gcc-toolchain
 - -DCMAKE_C_COMPILER_EXTERNAL_TOOLCHAIN
 - -DCMAKE_CXX_COMPILER_EXTERNAL_TOOLCHAIN
- Set --target
 - -DCMAKE_C_COMPILER_TARGET
 - -DCMAKE_CXX_COMPILER_TARGET
 - -DCMAKE_ASM_COMPILER_TARGET

Overriding the linker with CMake

- CMake will by default use the compiler driver
- For bare metal it can be useful to use arm-none-eabi-gcc instead of clang
 - Support for specs files
 - More mature linker script support
- Can use CMAKE_C_LINK_EXECUTABLE and CMAKE_CXX_LINK_EXECUTABLE to override.

```
set(CMAKE_C_LINK_EXECUTABLE "/path/to/arm-none-eabi-gcc <FLAGS>  
<CMAKE_C_LINK_FLAGS> <LINK_FLAGS> <OBJECTS> -o <TARGET>  
<LINK_LIBRARIES> -lc")
```

```
set(CMAKE_CXX_LINK_EXECUTABLE "/path/to/arm-none-eabi-g++ <FLAGS>  
<CMAKE_C_LINK_FLAGS> <LINK_FLAGS> <OBJECTS> -o <TARGET>  
<LINK_LIBRARIES> -lc")
```

Differences between GCC and Clang

- Assembler
- Predefined Macros
- LLD

GCC/Clang assembler differences

- Clang use its integrated assembler by default
 - Uses the same target selection options as Clang.
 - Is fussier on syntax than gas.
 - Does not permit coprocessor syntax for VFP/Neon in Armv7-a and above.
 - Unified Assembly Language (UAL) only.
 - Does not support `.altmacro`
 - Single pass so doesn't support some symbolic calculations.
- Can use `-fno-integrated-as` to use the GNU assembler
 - Uses heuristics to match target options to gas command line.

GCC/Clang general differences

- Clang will define many, but not all macros that GCC defines.
- In particular it claims to be GCC 4.2.1
 - `__GNUC__ 4`
 - `__GNUC_MINOR__ 2`
 - `__GNUC_PATCHLEVEL__ 1`
- There are Clang equivalents that ideally should be used instead
 - `__clang__ 1`
 - `__clang_major__ 9`
 - `__clang_patchlevel__ 0`
- Can get list via `clang -dM -E - < /dev/null`

LLD Limitations and differences

- LLD defaults to `-ztext` (no dynamic relocations in `.text` section)
- LLD has limited support for `-N` (`--omagic`), and no support for `-n` (`--nmagic`)
 - Can simulate with `-zmax-page-size=1`
- `PT_LOAD` program header generation not identical to `ld.bfd`
 - Can use `PHDRS` to select these manually.
- LLD tends to support subset of features that upstream Clang supports.
- LLD Linker Script support is not perfect
 - Some syntax differences “`()`” for address not allowed.
 - Chain of references
 - `alias_to__text = alias_to__text1;`
 - `alias_to__text1 = alias_to__text2;`
 - `alias_to__text2 = __text;`

Clang/LLVM toolchains today

- How to support Clang in your open source project.
- Assembling a Clang based toolchain today.

Supporting Clang in your open source project

- Rely on the cross-compilation options of the build system
 - Clang/LLVM can be cross-compiled using CMake.
 - Burden on the user to get the options right for their system.
 - Beware hidden host dependencies
 - DLLVM_TABLEGEN=/path/to/host/llvm-tablegen
- Supply the toolchain
 - Android NDK, Google Chrome are large enough projects to provide their own toolchains.
 - chromium/src/third_party/llvm-build/Release+Asserts/bin
 - Build systems pre-configured with all the paths and include directories.
- Customize Clang with via a new OS, Vendor, Environment
 - Conditionally alter existing ToolChain, for example Android modifies Linux.
 - Write a new ToolChain class that is selected via OS, Vendor, Environment
 - Linux, Fuchsia, FreeBSD ...

How could we assemble a Clang toolchain today?

- A linux toolchain akin to the arm-linux-gnueabi toolchain
 - Arrange the directory structure such that Clang can find it from the sysroot.
 - Set the default sysroot at build time, `DEFAULT_SYSROOT`.
 - Choose either `libc++` or `libstdc++` so that the includes don't clash.
 - Provide target compiler-rt libraries including sanitizers.
 - LLD as default linker.
- A bare metal embedded toolchain akin to arm-none-eabi toolchain
 - Partition libraries and includes by directory.
 - Provide compiler-rt builtins for the various m-class architectures.
 - Provide config files to replace specs-files and multilib
 - `armv6m.rdimon.nano.cfg`
 - Essentially provide the necessary command line options for the most common cases.
 - Config files could be separated into separate pieces so that users could construct their own.
 - Would probably need a GNU binutils, with `ld.bfd` as the default linker.
- In both cases some GNU libraries would need to be compiled with GCC.

Thank you

Join Linaro to accelerate deployment of your
Arm-based solutions through collaboration

contactus@linaro.org



96boards is a range of specifications with
boards and peripherals offering different
performance levels and features in a
standard footprint.



Linaro
connect

Bangkok 2019